

Deliverable 2 — Sprint Reflection

Patient Health Record Consent Management System • Agile Backlog Creation & Sprint Simulation in Jira

Name	SRN	Course	Jira Project
ABHAY S	PES1UG24CS013	Software Engineering - Lab 2	PHR

Sprint 1 committed 21 points across 4 stories; Sprint 2 committed 8 points across 2 stories. Project Epics total 29 points.

1. Did your estimations reflect the actual effort?

No story was actually coded during this simulation session, so there is no live developer hour metric to compare directly against. However, the honest evaluation is whether relative complexity was properly represented. On the complex items, US-04 (Provider Identity Verification Gateway, PHR-7) drew an 8-point estimate due to high uncertainty in integrating external national provider registries, mTLS certificate validation, and handling fallback timeouts. Conversely, US-02 (Instant Revocation, PHR-5) was estimated at 3 points because database flag updates and cache invalidation follow well-established architectural patterns. Thus, the Fibonacci progression (1, 2, 3, 5, 8) gave an accurate relative representation of risk and complexity.

2. Was your backlog well-prioritized?

Yes, delivery order followed the essential dependency hierarchy. A patient must have access to grant and revoke consent permissions (US-01, US-02), and clinic requests must be authenticated and validated against verified doctor records (US-03, US-04) before secondary systems can run. Therefore, Stories 1.1 through 2.2 correctly occupied Sprint 1. Real-time expiry alerts (US-05) and regulatory append-only audit logging (US-06) naturally consume the event stream produced by core transactions, making Sprint 2 their logical destination. This dependency-driven sequence eliminated blocking bottlenecks.

3. How did your simulated sprint align with your plan?

The simulated sprints executed in complete accordance with the plan. Sprint 1 committed 21 points across 4 stories, and Sprint 2 committed 8 points across 2 stories. All items transitioned systematically across To Do -> In Progress -> Done within their timeboxes, achieving a 100% completion rate with zero carried-over debt. In a real-world software delivery cycle, completing 100% of committed points in consecutive sprints might signal conservative estimation, but in this simulation it validated that scope was well-bounded and free of mid-sprint scope creep.

4. What insights did the burndown chart give about your team's capacity?

The burndown chart demonstrated an empirical velocity of 21 points for Sprint 1 and 8 points for Sprint 2. The stepped burndown pattern reflected that smaller stories (3 and 5 points) burn down incrementally, whereas large 8-point stories (US-04) keep the remaining work line flat until late in the iteration. This provides a valuable capacity insight: for production sprints, breaking down 8-point integration stories into smaller sub-tasks (e.g., separating mock schema validation from live registry mTLS connections) ensures a smoother burndown trajectory and earlier risk mitigation.